Uppsala University

Large Datasets For Scientific Applications

1TD268

Assignment A2

Author:
Sudarsan Bhargavan

June 9, 2018

1 A - Working with the RDD API

1.1 Question A.1

1. Read the English transcripts with Spark and count the number of lines.

   DataSet: Bulgarian dataset(europarl-v7.bg-en.en)
   Number of Lines: 406934

2. Do the same with the other language (so that you have a separate lineage of RDDs for each).

   No of Lines mentioned in the repository document : 406934

   WordCount Script

```
//WordCount.sh
  #!/bin/bash

  lineCount_en=‘/usr/bin/wc --lines ./europarl-v7.bg-en.en‘
  lineCount_bg=‘/usr/bin/wc --lines ./europarl-v7.bg-en.bg‘

  # Print Only The Number of Lines

  output_en=‘echo $lineCount_en |/usr/bin/awk ’{print $1}’‘
  output_bg=‘echo $lineCount_bg |/usr/bin/awk ’{print $1}’‘

  echo "Line Count of English Dataset: $output_en"
  echo
  echo "Line Count of Bulgarian Dataset: $output_bg"
```

   Output

```
  Line Count of English Dataset: 406934

  Line Count of Bulgarian Dataset: 406934
```

3. Verify that the line counts are the same for the two languages.

   Verifying the line counts for both the languages yields: [1]

   Bulgarian
   lines_en = sparkC.textFile("/home/ubuntu/DATA/europarl-v7.bg-en.en")
   lines_en.count()
   406934

   English
   lines_bg = sparkC.textFile("/home/ubuntu/DATA/europarl-v7.bg-en.bg")
   lines_bg.count()
   406934

4. Count the number of partitions.

   Trying to find the number of partitions yields:

   lines_en.getNumPartitioins()
   2

1.2 Question A.2

1. Inspect 100 entries from your RDD to verify your pre-processing.

   Inspecting 100 entries from RDD

   Output:

   For The Output Please See The File

   inspectedData.txt

2. Verify that the line counts still match after the pre-processing.

   Inspecting line counts after pre-processing still yields the same results for both the languages
   406934

1.3 Question A.3

1. Use Spark to compute the 10 most frequently according words in the English language corpus. Repeat for the other language.

   English

   Frequent Words List For English : [('the', 698563), ('of', 362452), ('to', 326291), ('and', 293700), ('in', 222084), ('a', 162764), ('is', 157336), ('that', 155812), ('for', 119429), ('I', 108253)]

   Bulgarian

   See File Referenced Below
   frequentsBulgarian.txt

2. Verify that your results are reasonable.

   After Translation Bulgarian —— English
   It was found that many matched with the frequent English words.
   Please See Matched Frequent Words

1.4 Question A.4

1. Do your translations seem reasonable?

   While manually comparing with google translate, the translation seemed reasonable.

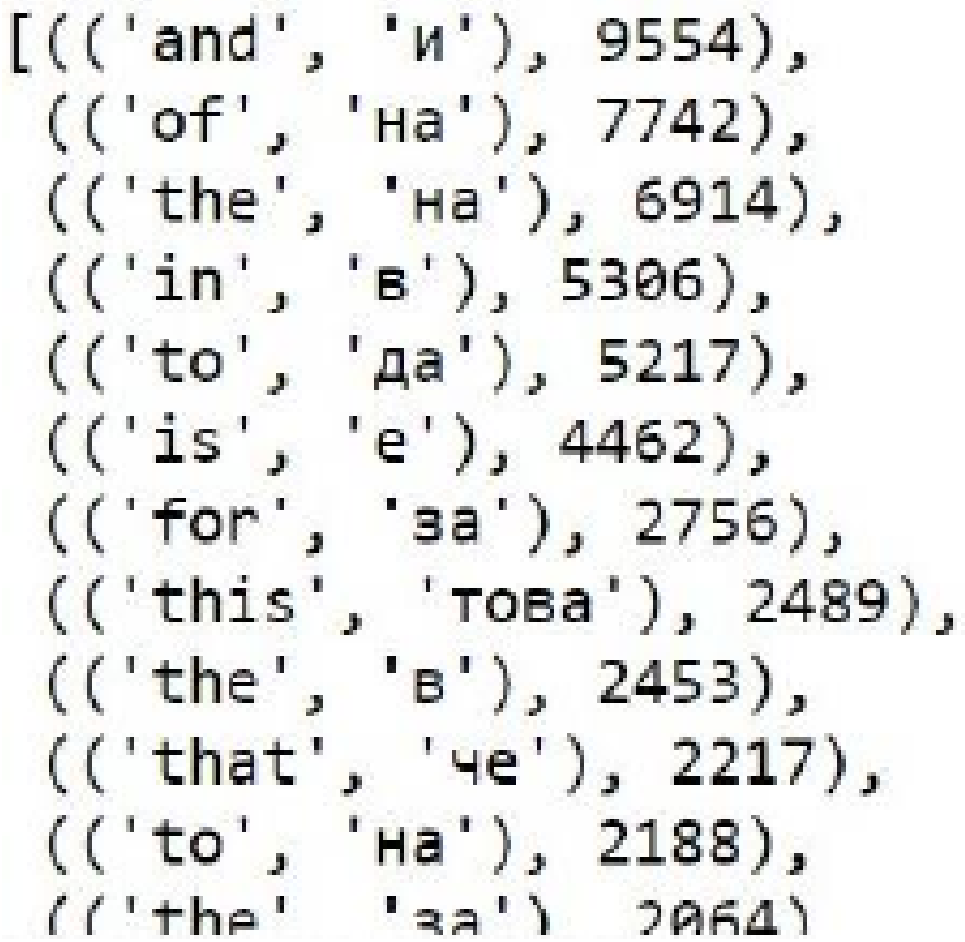

2 B - Working with DataFrames and SQL

2.1 Question B.1 - Analysis with DataFrames / SQL

1. Which organization has the largest gender pay gap? Which the least?

Largest Gender Pay Gap:

DiffMeanHourlyPercent	EmployerName
92.5	STOKECITYFOOTBALL...
88.4	BURNLEYFOOTBALL&A...
87.8	SWANSEACITYASSOCI...
87.7	MANCHESTERCITYFOO...
87.4	WESTHAMUNITEDFOOT...
87	WATFORDASSOCIATIO...
85.1	SUNDERLANDASSOCIA...
84.4	WESTBROMWICHALBIO...
84.4	SOUTHAMPTONFOOTBA...
84	CPFCLIMITED
83.3	NEWCASTLEUNITEDFO...
83	CHELSEAFOOTBALLCL...
83	MIDDLESBROUGHFOOT...
83	TottenhamHotspurF...
83	AFCBOURNEMOUTHLIM...
81.3	HARGREAVEHALELIMITED
79.6	THEARSENALFOOTBAL...
78	LEICESTERCITYFOOT...
78	SHEFFIELDWEDNESDA...
77.5	THELIVERPOOLFOOTB...

only showing top 20 rows

Least Gender Pay Gap:

DiffMeanHourlyPercent	EmployerName
0	ChoicesHousingAss...
0	BANBURYHEATHLIMITED
0	ErskineHospital
0	CINNAMONCARECOLLE...
0	ACCALIMITED
0	CMDRECRUITMENTLIM...
0	ANGELHUMANRESOURC...
0	COMFORTCALLLIMITED
0	AVENUECARESERVICE...
0	COOPERTOPCOLIMITED
0	24-7EMPLOYMENTSOL...
0	CRAIGTONFOODSLIMITED
0	BLUESAGENCYLIMITED
0	CYCLETRAININGUKLI...
0	BRAYBORNEFACILITI...
0	D.G.F.LIMITED
0	CAVITYDENTALSTAFF...
0	DALECARELIMITED
0	ACUMENLOGISTICSGR...
0	DAWSON&SANDERSONL...

only showing top 20 rows

2. What is the mean gender pay gap across all organization?

   Mean Gender Pay Gap:

```
+------------------------------------------------------------------+
|(sum(CAST(DiffMeanHourlyPercent AS DOUBLE)) / CAST(10491 AS DOUBLE))|
+------------------------------------------------------------------+
|                                                 14.298103136021377|
+------------------------------------------------------------------+
```

3. Export the results of B.1.2 to a CSV file. Inspect the output file to check it looks reasonable.

   Please See File
   csv.file

4. What proportion of organizations pay women more than men on average?

   Proportion of Organization That Pay Women More:

```
+--------+
|count(1)|
+--------+
|   10491|
+--------+

+-----------------------------------------------------+
|(CAST(count(1) AS DOUBLE) / CAST(10491 AS DOUBLE))|
+-----------------------------------------------------+
|                                  0.1167667524544848|
+-----------------------------------------------------+
```

2.2 Question B.2- Advanced DataFrames / SQL

1. Create a new column for the industry sector (for each company) using the SIC code:

   The **broadcast** and **join** variables were used to modify the **Data Frame**. Also as per the instructions given the **sic_codes** with value **-1** has been ignored.

   The **broadcast** variable is used to maintain a read-only cached data of the variable. Data has been joined as per the required conditions, with help of the **join** command. [2]

2. Compute the mean gender pay gap per sector.

   Mean Gender Pay Gap:

(sum(CAST(DiffMedianHourlyPercent AS DOUBLE)) / CAST(count(Industry) AS DOUBLE))	Industry
7.862613065326625	Wholesale_vehicles
8.059420289855073	Water_supply
9.660732984293192	Transportation
9.233670886075947	Support
11.603200000000003	Real_estate
9.785714285714286	Public_defense
14.778541953232475	Prof_sci
9.334634146341465	Other_service
13.76746812386155	Manufacturing
22.305303030303026	Insurance
17.868119266055047	Info_com
0.19999999999999998	Household
2.854654654654654	Health
6.022222222222222	Extraterritorial
15.651851851851852	Electricity
13.661538461538465	Education
23.853354632587862	Construction
6.594666666666666	Arts
3.743589743589742	Acc_food

(sum(CAST(DiffMeanHourlyPercent AS DOUBLE)) / CAST(count(Industry) AS DOUBLE))	Industry
14.909246231155768	Wholesale_vehicles
7.499999999999998	Water_supply
10.276178010471213	Transportation
11.227088607594942	Support
16.024709090909095	Real_estate
9.176190476190477	Public_defense
18.491334250343872	Prof_sci
12.46292682926829	Other_service
14.340364298724948	Manufacturing
26.281313131313123	Insurance
19.739228183486265	Info_com
3.133333333333333	Household
6.582132132132131	Health
9.944444444444445	Extraterritorial
14.785185185185187	Electricity
11.738219780219784	Education
21.771565495207675	Construction
21.06199999999999	Arts
7.8081318081318015	Acc_food
16.566666666666666	89

3. How does gender pay equality compare per sector? Compute some additional statistics.

   Calculating the mean values yields the following information :

   In some cases **women** were paid more than **mean**, but in most cases it was the other way around.

   While calculating median mean per sector the gender pay equality was **netural**.

(sum(CAST(DiffMedianHourlyPercent AS DOUBLE)) / CAST(count(Industry) AS DOUBLE))	Industry
7.862613065326625	Wholesale_vehicles
8.059420289855073	Water_supply
9.660732984293192	Transportation
9.233670886075947	Support
11.603200000000003	Real_estate
9.785714285714286	Public_defense
14.778541953232475	Prof_sci
9.334634146341465	Other_service
13.76746812386155	Manufacturing
22.305303830303026	Insurance
17.868119266055047	Info_com
0.19999999999999998	Household
2.854654654654654	Health
6.022222222222222	Extraterritorial
15.651851851851852	Electricity
13.661538461538465	Education
23.853354632587862	Construction
6.594666666666666	Arts
3.743589743589742	Acc_food

While calculating the mean per sector of **mean bonus pay**, it was found that **women** were paid more.

(sum(CAST(DiffMeanBonusPercent AS DOUBLE)) / CAST(count(Industry) AS DOUBLE))	Industry
-50.3013065326633	Wholesale_vehicles
9.556521739130435	Water_supply
13.2479057591623	Transportation
7.8084388185654054	Support
23.070399999999996	Real_estate
18.576190476190476	Public_defense
32.17345254470423	Prof_sci
17.52731707317074	Other_service
8.211256830601098	Manufacturing
46.9901515151515153	Insurance
36.38853211009177	Info_com
16.366666666666667	Household
-7.950900900900904	Health
24.166666666666668	Extraterritorial
27.09444444444434	Electricity
-15.838461538461546	Education
27.91246006389779	Construction
27.685	Arts
10.840659340659341	Acc_food
49.26666666666667	89

While calculating the median per sector of **mean bonus pay**, a lot of negative values were found, which according to the references provided, means that **women** were paid more. [3]

(sum(CAST(DiffMedianBonusPercent AS DOUBLE)) / CAST(count(Industry) AS DOUBLE))	Industry
-51.45336683417089	Wholesale_vehicles
-42.89999999999999	Water_supply
-35.23979057591623	Transportation
1.077130801687757	Support
12.811199999999996	Real_estate
18.89761904761905	Public_defense
13.30701513067399	Prof_sci
-4.016585365853658	Other_service
-45.433734861930814	Manufacturing
12.674242424242415	Insurance
-3.429128440366967	Info_com
10.0	Household
-3.0962462462462454	Health
12.266666666666667	Extraterritorial
22.235185185185188	Electricity
-13.298351648351646	Education
-3.457188498402587	Construction
-21.37633333333334	Arts
-15.151648351648364	Acc_food

3 C - Spark Clusters and Deployment

1. Modify a copy of your code from Section A, so that it runs on your cluster.

   Run Jobs - Cluster Mode
   In order to run a pyspark job in the cluster, the spark master url has to passed to the **SparkContext** method.

```
//PySparkJobClusterMode.py
    #!/usr/bin/env python3
    import pyspark as pys
       sparkC = pys.SparkContext("spark://localhost:7077")
```

2. Run your code first without and then with .cache() - and look under the storage tab in the web GUI for your application. What do you notice? Explain briefly what's going on.

   Cache When the cache method is used, the **RDD** caches a copy of the imported data, for further operations.
   When we omit the cache method, then the **RDD** waits for an event to get triggered, after which it loads the data. [4]

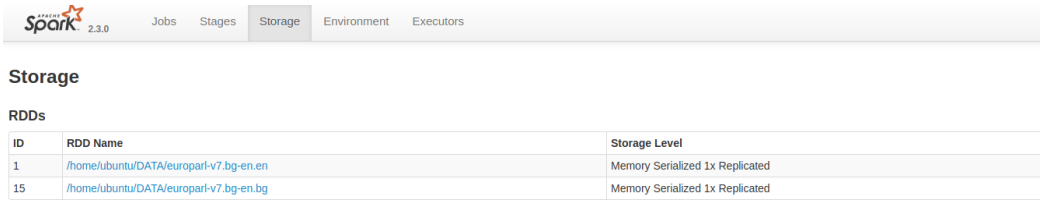

3. Use the Web GUI to explore your cluster and examine jobs, stages, and tasks. Create an example that requires a job with more than one stage. Explain, with reference to the Spark API methods you invoke in your code, why this is so.

   Multi-Stage Jobs
   A stage is a smaller set of tasks from a job. Stages can be parallelized if they are independent transformations are actions. [5]

   Here the **Task A3** has multiple-stages. But they cannot be parallelized because each stage is dependent on each other. **Stage Id: 0** represents the reduce operation. **Stage Id: 1** represents the sort operation.

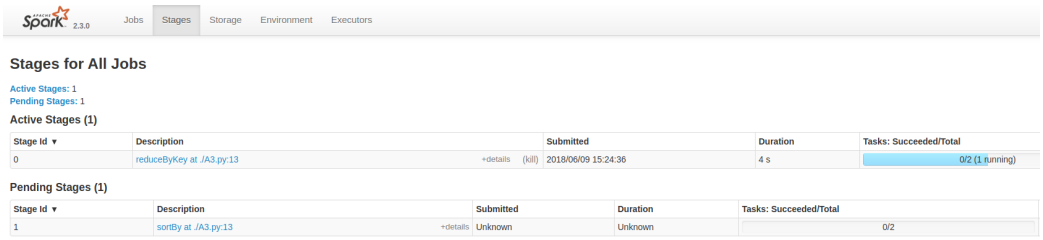

References

[1] statmt.org. (2018) European parliament proceedings parallel corpus 1996-2011. Accessed: 2018-06-09. [Online]. Available: http://www.statmt.org/europarl/

[2] learn4master.com. (2018) Pyspark broadcast variable example. Accessed: 2018-06-09. [Online]. Available: http://www.learn4master.com/big-data/spark/pyspark-broadcast-variable-example

[3] G. E. O. acas.org.uk, "Managing gender pay reporting," Tech. Rep., 2017-December.

[4] D. Darabos. (2018) (why) do we need to call cache or persist on a rdd. Accessed: 2018-06-09. [Online]. Available: https://stackoverflow.com/questions/28981359/why-do-we-need-to-call-cache-or-persist-on-a-rdd

[5] javadba. (2018) How are stages split into tasks in spark? Accessed: 2018-06-09. [Online]. Available: https://stackoverflow.com/questions/37528047/how-are-stages-split-into-tasks-in-spark